

SocialMediaDataAnalysis

March 17, 2025

1 Clean & Analyze Social Media

1.1 Introduction

Social media has become a ubiquitous part of modern life, with platforms such as Instagram, Twitter, and Facebook serving as essential communication channels. Social media data sets are vast and complex, making analysis a challenging task for businesses and researchers alike. In this project, we explore a simulated social media, for example Tweets, data set to understand trends in likes across different categories.

1.2 Prerequisites

To follow along with this project, you should have a basic understanding of Python programming and data analysis concepts. In addition, you may want to use the following packages in your Python environment:

- pandas
- Matplotlib
- ...

These packages should already be installed in Coursera's Jupyter Notebook environment, however if you'd like to install additional packages that are not included in this environment or are working off platform you can install additional packages using `!pip install packagename` within a notebook cell such as:

- `!pip install pandas`
- `!pip install matplotlib`

1.3 Project Scope

The objective of this project is to analyze tweets (or other social media data) and gain insights into user engagement. We will explore the data set using visualization techniques to understand the distribution of likes across different categories. Finally, we will analyze the data to draw conclusions about the most popular categories and the overall engagement on the platform.

1.4 Step 1: Importing Required Libraries

As the name suggests, the first step is to import all the necessary libraries that will be used in the project. In this case, we need pandas, numpy, matplotlib, seaborn, and random libraries.

Pandas is a library used for data manipulation and analysis. Numpy is a library used for numerical computations. Matplotlib is a library used for data visualization. Seaborn is a library used for statistical data visualization. Random is a library used to generate random numbers.

```
[1]: # your code here
import unicode

# Standard operational package imports.
import pandas as pd
import numpy as np

# Visualization package imports.
import matplotlib
import seaborn as sns

# Others
import calendar as cal
import re
import random

# Important imports for preprocessing, modeling, and evaluation.
from statsmodels.stats.outliers_influence \
    import variance_inflation_factor as smvif
import statsmodels.formula.api as smfapi
import statsmodels.api as smapi
import statsmodels.tools.tools as smtools
import statsmodels.stats.multicomp as smmulti
import sklearn.model_selection as sklmodslct
import sklearn.linear_model as skllinmod
import sklearn.metrics as sklmtcrs

[2]: # importing all my important data analysis functions
import data_analysis_functions
```

1.5 Generating Random Data

We are meant to generate Random Data for each field of data including categories such as “Fashion”, “Fitness”, “Music” “Culture”, “Politics”, “Family” and “Health” (which are ofcourse discrete variables)

Continuous variables must be data such as “No. Likes”, “No. Retweets” and “No. Views”

Datetime Variables such as “Date Tweeted” # we will look at bulk data in one year, since it is one company

```
[3]: # generating random dates in the year 2024
# intitilise dataframe
data_twitter = pd.DataFrame()
intended_sample_size = 650
```

```

# initialising random dates
start_date = pd.to_datetime('2024-01-01')
end_date = pd.to_datetime('2024-12-31')
days_in_year = (end_date-start_date).days # this corresponds to number of total
↳ tweets
data_twitter['Date'] = start_date + pd.to_timedelta\
    (np.random.randint(days_in_year,size=intended_sample_size), unit='d')
data_analysis_functions.df_head(data_twitter,10)

```

	Date
0	2024-07-28
1	2024-06-30
2	2024-09-02
3	2024-02-05
4	2024-10-26
5	2024-10-26
6	2024-06-28
7	2024-04-01
8	2024-06-30
9	2024-09-19

```

[4]: # categories
categories = ["Fashion", "Fitness", "Music", "Culture",\
    "Politics", "Family", "Health"]
data_twitter['Category'] = [random.choice(categories) for _ in
↳ range(intended_sample_size)]
data_analysis_functions.df_head(data_twitter,10)

```

	Date	Category
0	2024-07-28	Health
1	2024-06-30	Fitness
2	2024-09-02	Fashion
3	2024-02-05	Health
4	2024-10-26	Politics
5	2024-10-26	Music
6	2024-06-28	Music
7	2024-04-01	Fitness
8	2024-06-30	Fashion
9	2024-09-19	Fashion

```

[5]: # categories
realistic_likes_threshold = 2500
num_likes = np.random.
↳ randint(realistic_likes_threshold,size=intended_sample_size)
data_twitter['Num_of_Likes'] = [random.choice(num_likes) for _ in
↳ range(intended_sample_size)]
data_analysis_functions.df_head(data_twitter,10)

```

	Date	Category	Num_of_Likes
0	2024-07-28	Health	84
1	2024-06-30	Fitness	994
2	2024-09-02	Fashion	325
3	2024-02-05	Health	1446
4	2024-10-26	Politics	497
5	2024-10-26	Music	2199
6	2024-06-28	Music	799
7	2024-04-01	Fitness	506
8	2024-06-30	Fashion	1686
9	2024-09-19	Fashion	1210

```
[6]: # descriptive stats about our df
# print data types
data_analysis_functions.df_info_dtypes(data_twitter)
# descriptive summary
print(data_twitter.describe())
# counts of each category element
category_counts = data_analysis_functions.df_groupby_mask_operate(data_twitter,\
    'Category', 'Category', 0, '0', 'count')
print(category_counts)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 650 entries, 0 to 649
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Date             650 non-null   datetime64[ns]
1   Category         650 non-null   object
2   Num_of_Likes     650 non-null   int64
dtypes: datetime64[ns](1), int64(1), object(1)
memory usage: 15.4+ KB
None
```

	Date	Num_of_Likes
count	650	650.000000
mean	2024-07-05 02:12:55.384615424	1284.281538
min	2024-01-01 00:00:00	0.000000
25%	2024-04-03 00:00:00	618.750000
50%	2024-07-07 12:00:00	1341.500000
75%	2024-10-07 00:00:00	1934.750000
max	2024-12-30 00:00:00	2495.000000
std	NaN	755.921103

Category	count
Culture	92
Family	101

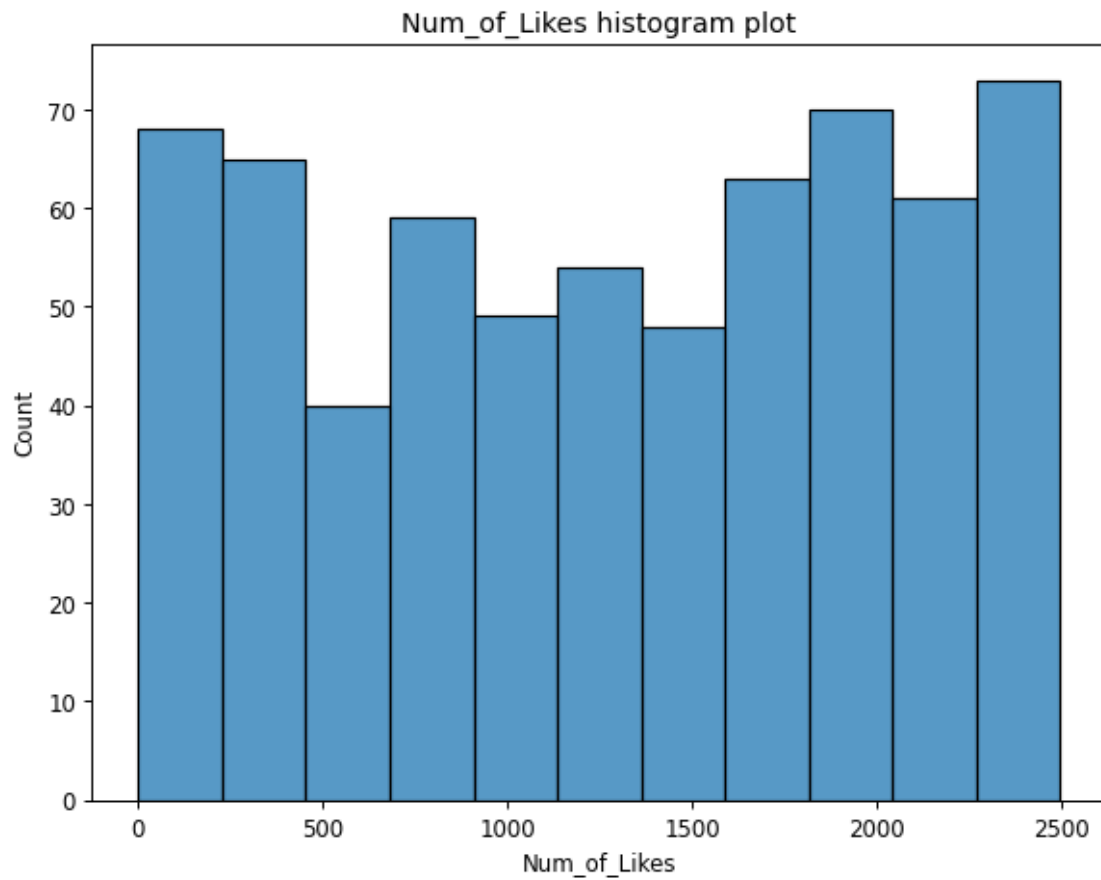
Fashion	102
Fitness	80
Health	96
Music	85
Politics	94

```
[7]: # removing all possible null data
data_twitter = data_twitter.dropna(axis=0).reset_index(drop=True)

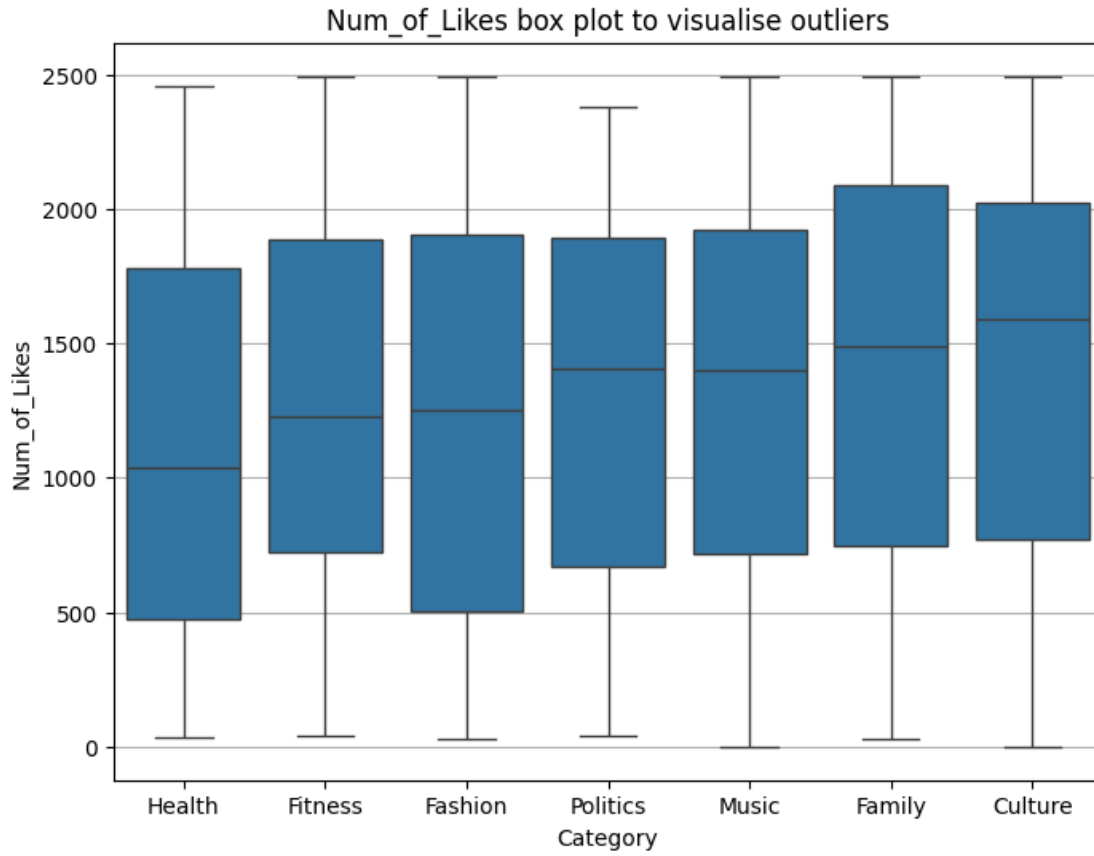
# convert dataframe date fields to datetime (already done)
data_analysis_functions.df_datetime_converter(data_twitter)
data_analysis_functions.df_head(data_twitter,10)
```

	Date	Category	Num_of_Likes
0	2024-07-28	Health	84
1	2024-06-30	Fitness	994
2	2024-09-02	Fashion	325
3	2024-02-05	Health	1446
4	2024-10-26	Politics	497
5	2024-10-26	Music	2199
6	2024-06-28	Music	799
7	2024-04-01	Fitness	506
8	2024-06-30	Fashion	1686
9	2024-09-19	Fashion	1210

```
[8]: # plot histogram of likes
data_analysis_functions.df_histplotter(data_twitter, "Num_of_Likes",2)
```



```
[9]: # boxplot of category
data_analysis_functions.df_boxplotter(data_twitter, "Category",
    ↪ "Num_of_Likes", 2)
```



```
[10]: # mean likes
mean_likes = np.round(data_twitter['Num_of_Likes'].agg(['mean']).values[0],2)
print("There are an average of {} Likes per tweet".format(mean_likes))
```

There are an average of 1284.28 Likes per tweet

```
[11]: # mean likes grouped by category
mean_likes_grouped = data_analysis_functions.
    ↪ df_groupby_mask_operate(data_twitter,\
                             'Category', 'Num_of_Likes', 0, '0', 'mean')
print(mean_likes_grouped)
```

	Num_of_Likes mean
Category	
Culture	1372.097826
Family	1406.702970
Fashion	1242.676471
Fitness	1244.950000
Health	1141.031250
Music	1305.976471

Politics 1272.095745

1.6 Early Inferences

We can see how there is no real shape to the distribution of Likes amongst categories, with most categories scoring a mean of 1100 likes. The box plot also shows the quartiles and outliers, which seem to be all similar amongst all likes categories.

We used here NumPy and random seeding and sampling, thus all the data in each category is independent of each

1.7 Advanced Hypothesis Testing using Regression Analysis

We can see perform a one-way ANOVA test on the Num_of_Likes using the Categories as a Categorical predictor. We do this by

```
[12]: # one way Linear Regression Analysis
data_analysis_functions.lr_ols_model(data_twitter, col_response="Num_of_Likes",\
                                     col_cont_predictors=[], col_cat_predictors=["Category"])
```

running simple linear regression model...

regressed variable: Num_of_Likes

continuous predictors: []

categorical predictors: []

```
[12]: {'Summary': <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  Num_of_Likes    R-squared:                  0.012
Model:                            OLS          Adj. R-squared:            0.003
Method:                 Least Squares          F-statistic:                1.330
Date:                  Mon, 17 Mar 2025         Prob (F-statistic):          0.241
Time:                   22:43:23                Log-Likelihood:           -5226.0
No. Observations:                650           AIC:                   1.047e+04
Df Residuals:                    643           BIC:                   1.050e+04
Df Model:                          6
Covariance Type:                  nonrobust
=====
coef    std err          t      P>|t|      [0.025
0.975]
-----
Intercept                1372.0978      78.690      17.437      0.000      1217.577
1526.619
C(Category) [T.Family]      34.6051     108.777       0.318      0.750     -178.996
248.207
C(Category) [T.Fashion]   -129.4214     108.523      -1.193      0.233     -342.523
```



```

83.681
C(Category) [T.Fitness]   -127.1478    115.382    -1.102    0.271    -353.720
99.424
C(Category) [T.Health]    -231.0666    110.119    -2.098    0.036    -447.304
-14.830
C(Category) [T.Music]     -66.1214    113.553    -0.582    0.561    -289.100
156.858
C(Category) [T.Politics] -100.0021    110.691    -0.903    0.367    -317.362
117.358

```

```

=====
Omnibus:                  624.880    Durbin-Watson:              2.139
Prob(Omnibus):            0.000    Jarque-Bera (JB):          43.135
Skew:                     -0.109    Prob(JB):                  4.30e-10
Kurtosis:                 1.757    Cond. No.                  7.91
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

""",

'Residuals': 0 -1057.031250

1 -250.950000

2 -917.676471

3 304.968750

4 -775.095745

...

645 73.904255

646 366.297030

647 159.323529

648 488.297030

649 -334.976471

Length: 650, dtype: float64,

'FittedValues': 0 1141.031250

1 1244.950000

2 1242.676471

3 1141.031250

4 1272.095745

...

645 1272.095745

646 1406.702970

647 1242.676471

648 1406.702970

649 1305.976471

Length: 650, dtype: float64}

1.8 Final Inferences

- 1) R-squared = 0.012, an R-squared this low essentially tells us the variance in Likes has nothing to do with Category of the tweets (0.012 is 1.2% of the variance in Likes is explained by the variance of Categories)
- 2) P-values, on the Health category with a P-value = 3.6%, has statistically significant effect on the number of tweets. We can confidently Infer that all other categories have no statistically significant impact on the Num_of_Likes per Tweets

We can see here that the Category of a tweet is likely to have next to no effect on the Number of Likes a tweet gets.

1.9 Functions I've Written for Data Analysis

Please check the `data_analysis_functions` for some bespoke data analysis functions that made this analysis 10 times easier.

ALL THE FUNCTIONS I WROTE MYSELF.